Group_Project_XL

install.packages("caret") library(caret) library(dplyr) install.packages("tidyverse") library(tidyverse) install.packages("cluster") library(cluster) install.packages("factoextra") library(factoextra) install.packages("cowplot") library(cowplot) library(ggplot2) install.packages("tidyr") library(tidyr) library(dplyr) install.packages("tidyverse") library(tidyverse) install.packages("cluster") library(cluster) library(readr) library(tidyr) install.packages("devtools") library(devtools) library(cluster) install.packages("fpc") library(fpc) library(readr) library(dplyr) library(ggplot2) install.packages("ggcorrplot") library(ggcorrplot) library(tidyr) library(fastDummies) library(caret) install.packages("VIM") library(VIM) library(readr) library(tidyverse) library(caret) library(pROC) library(ggcorrplot) library(gmodels) library(rpart)

Churn_Train <- read_csv("Churn_Train.csv") Churn_Data <- read_csv("Churn_Train.csv")

Inspecting data

head(Churn_Data)

Examining the dataset

glimpse(Churn_Data)

Summary statistics of dataset

summary(Churn_Data)

From glimpse we can see that, Some of the character variables can be converted into factors, So Converting character variables to factors.

Churn_Data <- Churn_Data %>% mutate_if(is.character, as.factor)

Checking NULL values in the dataset at column level.

colSums(is.na(Churn_Data))

imputation of missing values - median imputation technique

imputation_model <- preProcess(Churn_Data %>% select_if(is.numeric),method = "medianImpute") data <- predict(imputation_model, Churn_Data %>% select_if(is.numeric))

Churn_Data <- Churn_Data %>% select(setdiff(names(Churn_Data), names(data))) %>% cbind(data)

Box plot - to detect the outliers

Churn_Data %>% select_if(is.numeric) %>% mutate_all(scale) %>% gather("features","values") %>% na.omit() %>% ggplot(aes(x = features, y = values)) + geom_boxplot(show.legend = FALSE) + stat_summary(fun = mean, geom = "point", pch = 1) + # Add average to the boxplot scale_y_continuous(name = "Variable values", minor_breaks = NULL) + scale_fill_brewer(palette = "Set1") + coord_flip() + theme_minimal() + labs(x = "Variable names") + ggtitle(label = "Distribution of numeric variables in Churn dataset")

Visualizing distribution of Churn categorical variable.

ggplot(Churn_Data, aes(x=churn, y=..prop..,group = 1)) + geom_bar(fill="light blue") + theme_classic() + geom_text(aes(label=round(..prop..,2)),stat = "count", position = position_stack(vjust=0.5)) + labs(y = 'Proportion', title = "Proportion of churn") + scale_x_discrete(labels = c("No","Yes"))

finding correlation between variables

Churn_Data_cor <- round(cor(Churn_Data %>% select_if(is.numeric)), 1)

ggcorrplot(Churn_Data_cor, title = "Correlation", type = "lower") + theme(plot.title = element_text(hjust = 0.5), axis.text.x = element_text(angle = 90))

Total minutes and total charge for the day, evening, night, and international are strongly linked, we can deduce.

Churn_Data <- Churn_Data %>% select(-state, -churn) %>% fastDummies::dummy_cols(., remove_selected_columns = TRUE) %>% mutate(state = Churn_Data$state, churn = $Churn_Data$churn)

Pre-Processing of data

Splitting dataset into training (80%) and validation (20%) sets

set.seed(12) index <- createDataPartition(Churn_Data$churn, p=0.8, list=FALSE) Churn_Data_train_df <- Churn_Data[index,] Churn_Data_test_df <- Churn_Data[-index,]

scaling the data

scaling <- preProcess(Churn_Data_train_df %>% select_if(is.numeric), method = c("center", "scale"))

Churn_Data_train_norm <- predict(scaling, Churn_Data_train_df %>% select_if(is.numeric))

Churn_Data_test_norm <- predict(scaling, Churn_Data_test_df %>% select_if(is.numeric))

Churn_Data_train_norm$churn <- $-Churn_Data_train_df$churn Churn_Data_test_norm$churn <- $-Churn_Data_test_df$churn

Model Construction

Model_1 <- glm(churn ~ ., data = Churn_Data_train_norm , family= "binomial")

summary(Model_1)

Predict values using based on Model_1.

pred_probs <- predict(object = Model_1,Churn_Data_test_norm, type = "response")

Assigning labels based on probability prediction

Model_Pre_lables <- as.factor(ifelse(pred_probs>0.6 ,"yes","no"))

Performance Metrics

Confusion matrix for significant variable model.

confusionMatrix(Model_Pre_lables,Churn_Data_test_norm$churn)

AUC of the churn model

roc(Churn_Data_test_df$churn, pred_probs)

plot.roc(roc(Churn_Data_test_df$churn, pred_probs))

Applying the model to the Customers to Predict data file

Load the data file

load("C:/Users/xlamo/Desktop/XanLamoreux/Group Project/Customers_To_Predict.RData")

creating a copy to work with

customer_predict <- Customers_To_Predict

removing the state column as it is not necessary

customer_predict <- customer_predict %>% select(-state) %>% fastDummies::dummy_cols(., remove_selected_columns = TRUE)

Transformation for scaling the data (Z score transformation)

customer_predict <- as.data.frame(scale(customer_predict))

#predicting the model with the test data

predict_labels <- predict(object=Model_1,customer_predict,type="response")

applies the probability ratio if under 60% customer will not churn

Model_Pre_lables_2 <- as.factor(ifelse(predict_labels>0.6 ,"yes","no"))

adding chrun column and attaching the predictor from the model

Customers_To_Predict <- Customers_To_Predict %>% mutate(churn=Model_Pre_lables_2)

visual of the results which shows that 267 will churn

table(Customers_To_Predict$churn)

View(Customers_To_Predict)

The Customers_To_Predict file can be exported as the final results of our model